

Cracking the Black Box

Portfolio Replica Strategy — BC3 final deliverable

PoliMI Fintech — Business Case 3

Problem and data

Goal. Replicate a hidden “Monster Index” using only liquid Bloomberg futures, subject to a regulatory VaR cap.

$$r_t^{\text{target}} = 0.50 r_t^{\text{HFRX}} + 0.25 r_t^{\text{MSCI World}} + 0.25 r_t^{\text{Global Bond}}$$


Universe (11 futures).

- Rates: RX1, TY1, DU1, TU2
- Equity: ES1, NQ1, VG1, TP1, LLL1
- Commodity: GC1, CO1

Constraints.

- Sample: weekly, Oct 2007 – Apr 2021 (704 obs).
- Gaussian VaR cap (1%, 1-month) $\leq 8\%$.
- Gross-exposure cap: $\sum_j |w_j| \leq L$.
- Transaction cost: $\tau = 5$ bps on L1 turnover.

Pipeline overview — one harness, four model families



Backtest methodology — Part I

Rolling walk-forward. At every rebalance week t :



1. **Fit** on $X[t-104 : t]$, $y[t-104 : t]$ (no look-ahead — $X[t]$, $y[t]$ never seen).
2. **Project** onto the feasible set: gross cap $\sum_j |w_j| \leq L$, then iterative shrinkage until $\text{VaR}(w) \leq 8\%$.
3. **Trade** $\Delta w_t = w^{\text{new}} - w^{\text{old}}$; **cost** $c_t = \tau \cdot \sum_j |\Delta w_{t,j}|$.
4. **Realise** $r_t^{\text{rep}} = w_t^\top r_t$; on non-rebalance weeks carry w_{t-1} forward.

Contract metrics.

$$\text{TE} = \sqrt{52 \text{Var}(r_t^{\text{rep}} - y_t)}, \quad \text{IR} = \frac{52 \mathbb{E}[r_t^{\text{rep}} - y_t]}{\text{TE}}, \quad \text{net_IR} = \frac{52 \mathbb{E}[r_t^{\text{rep}} - c_t - y_t]}{\text{net_TE}}$$

VaR fallback. Jarque-Bera on the target rejects normality ($p \approx 0$); Gaussian VaR (4.14%) understates the historical 1%-quantile (5.96%) — the harness can fall back to a historical VaR projector.

Part II — classical linear benchmark



Sweep. {OLS, Ridge, Lasso, EN, Huber} $\times$ {all-11, top-5} — single harness: 104-week window, $\Delta t = 4$, scale-only, GE cap 2, VaR cap 20%.

Top- K . Keep the K futures with highest $|\text{corr}(r_j, y)|$ inside the training window — no

Reading.

- Ridge / EN collapse to OLS at $\alpha < 10^{-3}$.
- Top-5 pruning saves $\sim 4\times$ on average GE for a small TE premium.
- Sparsity from Lasso buys turnover savings but loses on tracking precision.


Part III — portfolio constraints (post-hoc projection)

Why projection matters. Gaussian VaR is rejected by Jarque-Bera; the historical 1%-quantile is fatter than $\mathcal{N}$. The harness projects w onto a *historical* VaR feasible set.

Six base configurations (B1–B6):

$$\Delta t \in \{1, 4, 13\} \text{ wk} \times L \in \{1.0, 2.0\}$$

Crisis-window evaluation. 2008-09 $\rightarrow$ 2009-03 (GFC) and 2020-02 $\rightarrow$ 2020-04 (COVID). The COVID drawdown is the binding test: every config beats the target on max-drawdown, but only the quarterly cadences keep TE bounded.



Part III — predict-then-optimize (convex layer)

Two-stage rebalance.

1. **Alpha.** Fit a classical linear model on (X_{tr}, y_{tr}) ; take $\mu = \text{coef}_-$.
2. **Risk.** Ledoit–Wolf shrunk sample covariance Σ (annualised).
3. **Optimise** (CVXPY):

$$\max_w \mu^\top w - \frac{1}{2} \lambda w^\top \Sigma w - \tau \|w - w_{t-1}\|_1 \quad \text{s.t.} \quad \|w\|_1 \leq L, \quad |w_j| \leq 0.5$$

Sanity check. With $\lambda = 2$, slack constraints, no LW shrinkage, the convex layer recovers OLS to within solver tolerance.

Findings.

- Longest rolling window (3y) and modest Ridge dominate net IR on the leaderboard.
- Huber is competitive only in the 2008/2020 windows.
- Lasso's sparsity buys lower turnover but loses on TE.

Part IV — Kalman filter / state-space

Random-walk weights model.

$$x_t = x_{t-1} + B u_t, \quad u_t \sim \mathcal{N}(0, I_{11})$$


$$y_t = r_t^\top x_t + D \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, 1)$$

$B = \sigma_w I$, $D = \hat{\sigma}_y$; Ridge warm-start on the first 52 weeks.

Variants:

- Static σ_w grid $\{10^{-4}, 10^{-3}, 10^{-2}\}$.
- EM-tuned (B, D) , capped at 3 iterations.
- Double-adaptive: σ_w and D rescale with rolling vol.


Leakage rule. Filter outputs $\hat{x}_{t|t}$ use y_t , so we *shift weights by one week* before evaluation.



Adaptive Kalman replica vs target.

Part V — deep learning weight generator (architecture)

End-to-end. The network outputs the *weights*, not the returns: $w_t = f_\theta(\phi_t)$, $r_t^{\text{rep}} = w_t^\top r_t$.



Loss — forward-window MSE + turnover + budget:


$$\mathcal{L}(\theta) = \underbrace{\text{MSE}_{s \in [t, t+H-1]}(w_t^\top r_s - y_s)}_{H=12} + \lambda_1 \|\Delta w\|_1 + \lambda_2 \|\Delta w\|_2^2 + \gamma (1^\top w_t - 1)^2$$

Part VI — consolidated leaderboard (top of 30 rows)


Model	net_IR	IR	ρ	TE	GE	$\bar{T}$
NN $\Delta t = 4$	+0.170	+0.177	0.841	2.96%	1.02	0.008
NN $\Delta t = 1$	+0.135	+0.143	0.843	2.93%	1.03	0.010
M3-PO Huber $\delta = 1.35$ all-11	+0.017	+0.032	0.856	2.93%	1.28	0.017
M3-PO OLS all-11	+0.017	+0.032	0.853	2.95%	1.38	0.018
M2 Huber all-11	-0.086	-0.022	0.844	3.03%	1.59	0.075
Kalman (basic, $\sigma_w = 10^{-3}$)	-0.141	-0.137	0.868	3.04%	0.59	0.005
M2 Lasso / EN top-5	-0.185	-0.173	0.828	3.18%	0.37	0.014
M3-Cstr B1 (weekly / 100% cap)	-0.468	-0.461	0.765	3.91%	0.19	0.011

Headline. NN $\Delta t = 4$ tops the leaderboard on net IR with $\rho = 0.84$ at the lowest gross exposure and turnover of the positive-IR group. The gap to the next cluster (PO + Huber / OLS, $\rho \approx 0.85$) is narrow; M3-Cstr (post-hoc EN at 1y window) loses the most — the constraint sweep without the right model class cannot rescue tracking quality.


Best model spotlight — cumulative & weights



NN $\Delta t = 4$ vs target, OOS 2011-11 $\rightarrow$ 2021-02.



Rolling 26-week correlation; overall $\rho = 0.84$.



Asset weights over time (RdBu = long/short).

Reading the heatmap.

- Persistent long in rates (RX1, TY1) carries the bond leg.
- Equity loadings (ES1, NQ1, VG1) rotate with the regime.
- Sign-flips are mild — the network respects the turnover penalty.

Findings

- **Methodology before models.** Part I's harness fixes rolling-window discipline, cost accounting, and the VaR / GE projection *once*; every Part II–V model plugs into the same evaluation surface and lands in the same master table.
- **Cadence sets the cost/fit trade-off.** Weekly rebalancing maximises tracking but pays the heaviest 5 bps drag; quarterly minimises drawdown but loses precision. The B1–B6 sweep makes the trade-off auditable.
- **Long window beats reactive fitting.** 3-year rolling dominates 1-year on net IR across every linear model — weekly data is too noisy to adapt aggressively.
- **Kalman EM matches the static $\sigma_w = 10^{-3}$ grid point.** The regime-switching variant narrows TE in 2008 at the cost of turnover.
- **The MLP wins net IR.** NN_dt4 delivers $\rho = 0.84$ at $\overline{GE} \approx 1.0$ and $\bar{T} \approx 0.8\%$; without $\lambda_1 \|\Delta w\|_1$, NN variants over-trade and lose net IR.
- **Test vs. val gap is regime, not overfit.** $\text{train} \approx \text{val} \ll \text{test}$ across λ_1 values — the test fold crosses COVID; the model is honest about this.