

Detecting Abnormal Markets

Early Warning Systems with Deep Anomaly Detection

Peng Rao, Maxime d'Aviau de Ternay, Qiao Zhaochen, Shourya Vashistha, Arnau Nomen

May 24, 2026

Politecnico di Milano

Why this matters

Markets crash often enough to matter – 1929, 2000, 2008, 2020 ...

Missing the 10 worst days vs. buy-and-hold S&P 500 (1927–2015):

Strategy	Growth of \$1
Total cumulative return	\$115.40
Miss the 10 worst days	\$362.01
Miss the 10 best days	\$38.28

Tail events cluster

Frequency and cross-asset correlations of simultaneous heavy losses keep rising. Diversification fails exactly when you need it most.

The problem: risk-on → risk-off

Risk-on

Investors are confident. They bid up equities, credit, EM debt, commodities.

Volatility is low. Spreads compress. Capital flows toward growth.

Risk-off

Investors are scared. They flee to cash, short-term govies, gold, CHF, JPY.

Volatility spikes. Correlations surge. This is systemic risk.

- We do not need to *forecast* the crash – we need to *nowcast* the transition.
- Goal: detect the green-to-red turn as early as possible.

Solution: Early Warning Systems

- Treat **normal markets** as the inlier distribution and **stress regimes** as anomalies.
- Train novelty detectors on the normal samples only; raise a defensive flag when the anomaly score crosses a threshold.
- Why this framing beats outright crash prediction:
 - No need to label every crisis type ex ante.
 - The model scores the latest weekly observation in real time.
 - It generalises – a 2008-style stress and a 2020-style stress both look “not normal”.

Deployment

A defensive flag rotates the portfolio into cash, short-term govies, gold, or other low-beta assets. Typical user: CIO, risk officer, wealth manager.

Bloomberg weekly panel

- Key equity & bond indices (IG/HY, inflation-linked, mortgages)
- Short / medium / long interest rates
- Major FX, commodities, VIX
- Economic Surprise, Baltic Dry
- Binary label: *normal* vs. *abnormal*

Same model, two business cases

- **Risk management** – reduce drawdown
(*this work*)
- Quant strategy – generate alpha

That choice drives data treatment, threshold tuning, and the evaluation metric we report.

Three deep novelty detectors

All three are trained on **normal observations only**; high score $\Rightarrow$ defensive flag.

- **GAN** – generator learns the normal manifold; two scoring heads
 - Discriminator realism: $s(x) = 1 - D_{\text{prob}}(x)$
 - Reconstruction in latent space: $s(x) = \min_z \|x - G(z)\|^2$
- **VAE** – reconstruction error of a Gaussian-regularised latent space; KL-regularised encoder gives a smoother score.
- **LSTM autoencoder** – reconstructs a rolling window of standardised features, adding short-term temporal context.

Threshold choice and evaluation

Threshold tuned on cross-validation

Maximise F_β with $\beta = 1.5$, subject to a precision floor of $\min(0.7, \text{base rate} + 0.25)$.

Headline metrics for the risk officer:

- **Recall** – fraction of real crises the detector would catch.
- **PR-AUC** – threshold-free summary that stays honest under heavy class imbalance.
- **False-negative rate** – a missed crisis is the expensive failure mode.

Precision, F1, ROC-AUC are reported for completeness but are not the primary lens: a defensive false alarm costs basis points; a missed crash costs a drawdown.

Takeaway

- Deep novelty detection turns a labelled crisis history into a real-time risk dial.
- The risk-managed equity curve trims drawdown vs. buy-and-hold, with mild return drag in calm regimes – the trade-off the risk-management user wants.
- Recall-weighted thresholds keep the detector useful when crises are rare.
- Next steps: ensemble the three scorers, add macro / credit-spread features, stress-test thresholds across regimes.

Detect the transition. Act on it. Reduce drawdown.